

Deep Learning for Image Recognition

Machine Learning Engineer Nanodegree

Capstone Project

Brian Lester

blester125@gmail.com

Abstract

Deep Learning has become one of the most popular specializations in Machine Learning due to the massive increase in both the amount of data available and the amount of computing power available. With increases in both of these the size and possible depth of neural networks has been increase, which has allowed for them to be applied to various new fields, especially image processing. Image processing is often done using Convolutional networks that reduce the size of images while increasing their depth. This paper applies a deep convolutional neural network to the problem of classifying multiple digits in natural scenes. The depth of this network allows for one network to classify each digit rather than traditional approach of localize, segment, and classifying each digit individually.

1. Definition

1.1. Project Overview

This project tackles multidigit sequence recognition in real world scenes. This has many applications, the most obvious being recognizing addresses from GPS tagged photos. This would allow for the creation of maps with addresses automatically using data collected (like the kind captured by the Google streetview car). This also has

other applications. As self driving cars become more popular the cars need a way to determine the speed limit. Without extensive Vehicle to Infrastructure communication a car needs to know the speed limit in an area, this can be done from map data or by reading the Speed Limit sign. This is especially in the case of road construction where map data may not be up to date. Multi digit recognition can also be used to automatically scan passports and the like. Clearly this technology can be applied in many areas.

One of the first applications of Convolutional Neural Networks was by LeCun in 1989 [5]. These networks were used for digit recognition so that mail sorting machines could recognize hand written digits that appear in the zip codes of mail. This allowed for automatic mail sorting/routing machines to be created. It is fitting that Neural Networks are still used for digit recognition. However current applications (even automatic mail sorters) need to be able to recognize more than one digit at a time (Zip Codes are 5 digits long after all). Using the old neural networks the problem had to be broken into several different parts. This included digit detection where digit sequences are found in the image, then this detected area is segmented into probable digits. These segments are then analyzed and individual digits are recognized. This multistep process meant the programmers have to write code for each step of the process. This

has the unfortunate consequence of taking up programmer time. It would be much easier to a larger neural network to do the whole process. The use of an end-to-end network to recognize all the digits in an image at the same time was first used by Goodfellow *et al.* [3]. Datasets for this project are the “Street View Housing Numbers” (SVHN) dataset that can be found here <http://ufldl.stanford.edu/housenumbers/>. This dataset includes 73257 training digits, 26032 training digits and 531131 extra “somewhat less difficult samples” to use as additional training examples. These images are RGB images with labels and bounding boxes included with each image.

1.2. Problem Statement

The problem addressed in this paper is the problem of finding and reporting multiple digits in a picture of a natural scene. This problem differs from the normal approach used to classify multiple digits in that it uses a single deep neural network to identify all the digits in the image rather than using a multiple step process of localizing the digits, segmenting them, and then classifying each digit individually. This deep network will learn to do these steps automatically without human help. This allows for the computer to optimize and segment the process in ways that are non-obvious to a human programmer.

This is a supervised learning problem and we will be solving it using a deep convolutional neural network. We will process the image with the convolutional network to create a much deeper feature vector. This new feature vector will then be passed to 6 fully connected layers (each of the 6 are two layers deep) that will output the length of the sequence in the image and the first 5 digits in the image. This end-to-end strategy is based on the findings of the Goodfellow *et al.* paper [3].

This deep neural network will output six softmax classifiers (this includes the probability of each possible value for that classifier). The first classifier will output the length of the sequence of

digits in the scene. The remaining five classifiers will output the value of each digit in the sequence (with 10 meaning that the digit is not present in the sequence). By outputting the argmax (index of the largest value of the softmax classifier) and taking the last five numbers without including the tens the resulting numbers will be the sequence of the digits in the numbers.

1.3. Metrics

That metric that will be used to measure the accuracy of the Deep Convolutional Neural Network is the proportion of correct classifications. For any given image there are six labels. The first is the length of the sequence of digits (1 through 6 where 6 means more than 5 digits). The next 5 labels are the digits in the image from 0 to 10 where 10 means that that particular digit is not in the image. This metric is sufficient for this project because recognizing multiple digits with a single neural network is the point of the project and how many classifications the network makes that are correct is a good metric for how well the network is performing. This is the accuracy value that was used in training and for evaluating the validation set. This metric is good for evaluating the model for this project.

If this network was used in an application where accuracy is even more important (for example creating addresses in automatically generated mapping application where a wrong address is a huge problem for a user) then a secondary accuracy metric could be used where the entire classification must be correct for the test sample to be correct. For example if the label is 137 and the model outputs 17 it is wrong not 66.6% correct. This network will not be used in this sort of high pressure situation so the metric that is simply the proportion of the correct classifications is sufficient for our purposes of seeing what tweaks to the model cause better performance.

2. Analysis

2.1. Data Exploration

The dataset comes from the Street View House Numbers dataset from Stanford. This dataset can be found here <http://ufldl.stanford.edu/housenumbers/>. This data is split into three datasets, the “Training” dataset with 33,402 images, the “Test” dataset with 13,068 images that are used to evaluate the performance of the network, and the “Extra” dataset that contains 202,353 images. The extra dataset images are considered “easier” examples than the training dataset. These images are called “easier” both by the website where the datasets can be downloaded and in a paper by Lecun *et al.* [6].

The dataset is a large collection of images taken from the Google streetview. The images are similar to MNIST dataset of hand written digits in that the images are small cropped digit images. However the difference is that the images from SVHN are taken from real-world, natural, scenes. The SVHN images also contain multiple digits. Each digit ranges from 0 to 9 and there are between 1 and 6 digits in an image. These images are RGB images of various sizes with three color channels (a Red, Green, and Blue channel). There are no other features outside of the values of the image pixels. The dataset also include labels for the digits. The dataset also contain the information to draw bounding boxes around each digit (they are given by the x and y coordinates and the width and height of the box). These boxes were used in data preprocessing to help center and crop digits from the images. The boxes are not considered input to the neural network. The pixels of the images are the only features that are feed to the Neural Network.

Table 1 shows some statistics about the datasets. Do to the order that the image processing is done (the images from the training and extra datasets are processed then split into training and validation sets) most of the data is not available for the final training set or the validation set. The

Length	1	2	3	4	5	6
Train	5,137	18,130	8,691	1,434	9	1
Extra	9,385	71,726	106,789	14,338	115	0
Test	2,483	8,356	2,081	146	2	0

Table 2. This table shows the frequency of various lengths of digits that are found in the Training dataset.

height is the max height of an image in that particular dataset, the width is the max width of an image in that dataset. The mean and standard deviation are the mean and standard deviation of the pixel values in each dataset. The table shows that all the dataset are approximately the same outside of the size. The training and validation sets do not have this sort of information because the images are processed (which includes normalization) before they are split into the two later dataset. This normalization means that the mean and standard deviation for these datasets are zero and 1 respectively.

Table 2 shows the length of various sequences in the datasets. These numbers are collected across all of the datasets. This table shows that there are a similar number of most of the lengths. There are very few images with five digits and only one that has six digits. This sequence of length 6 could be considered an outlier; however, it was not removed from out dataset because in the real world some addresses or images will have 6 or more digits and the model should be able to handle it. The table shows that the Training set has a mode of length 2, the extra set of length 3, and the test set of length 2.

The datasets are pretty simple, just a bunch of images that are all quite similar to one another.

2.2. Exploratory Visualization

The dataset is just pictures and the features are just pixel values. There are not various features that can be plotted against each other to find features that are correlated or related. Lacking this sort of visualization means feature cannot be found that could be removed. Instead of plotting features examples from the training dataset have

Dataset	Size	Max Height	Max Width	Mean	Std. Dev.
A.) Train	33,402	501	876	139.22	59.67
B.) Extra	202,353	415	668	135.94	61.62
C.) Test	13,068	516	1083	133.72	66.11
D.) Train	230,071	N/A	N/A	N/A	N/A
E.) Valid	5,684	N/A	N/A	N/A	N/A

Table 1. This table details information about the datasets. These datasets include A.) The training set that was split to create the D.) Train dataset and E.) Valid dataset, B.) The extra dataset that is used to create the training and validation datasets, C.) The test set that is used to evaluate the model, D.) The training set that is used for training the model (created from the train and extra datasets), and E.) The validation set that is used to tune hyperparameters of the model. Size is the number of images in each dataset. Max height is the height of the tallest image in each dataset. Max width is the width of the widest image in each dataset. Mean is the mean pixel value of the images in the dataset. Std. Dev. is the standard deviation of the pixel values in the images. *Note:* The final Training dataset and Validation sets have little information about them because they are created after the images from the Train and Extra datasets are processed and therefore have been normalized so they have no meaningful mean or standard deviation.




Figure 1. An example image from the Training dataset. Bounding boxes for each digit has been added to the image.

been added. Figure 1 is an image from the training dataset that has labels of length: 3 and values: 1, 0, 2, 10, 10. This means that the image contains the digit sequence 102 as can clearly be seen. This example image also includes the bounding boxes that surround each digit. These could be used to crop individual digits to train a single digit classifier but are instead used to help center and crop images so that all the digits are visible.

Figure 2 contains an example of an image that could be considered an outlier. This is the only digit in the Training set that contains a sequence of length six. This is not removed from the dataset however because in real world use of this network capturing sequences that include more than five




Figure 2. The only example in the training dataset that has six digits.

digits seems fairly common. In the case of a mapping application that is trying to find addresses that are up to length five it is important to be able to find sequences that are longer than five digits and disregard them rather than creating two outputs that overlap or some other unexpected prediction.

2.3. Algorithms and Techniques

The following algorithms will be used to implement a deep convolutional neural network to identify multidigit sequences in real world scenes.

- **Logistic Regression:** Using a matrix and a vector you can transform an input of size i into an output of size o by multiplying the input by an $[i, o]$ matrix then adding a vector of size o to it. By adjusting the value of

the matrix an output can be created that depends on different parts of the input with different weights based on the matrix. Logistic Regression is the defacto standard for neural networks which is why it is used here.

- **Rectified Linear Unit (ReLU) Activation:** Rectifier activation is used to introduce non-linearity to the network. When passed through a ReLU rectifier input is set to zero when below a threshold (the default tensorflow threshold is used). ReLU is used because it causes sparse activation (only some of the inputs are active) and it is efficient to compute.
- **Convolution:** Using a small matrix called a kernel that is size x by x by depth of the input (depth1) by depth of output (depth2) the kernel is moved over the input (the amount the kernel is moved is called the stride) that is of depth1 and transforms the input into an output that is slightly smaller but is now depth2. There are two forms of padding that can be done, valid padding where the kernel never leaves the image and same padding where the image is padded with zeros. This is a form of weight sharing where the same weights are applied to various parts of the image so that the network does not have to deal with where in the input the important features are. As the kernel moves across the image it performs logistic regression on each piece of the input.
- **Max Pooling:** Small sections of the image are grouped together and the largest value is used. This reduces the size of the input while retaining the depth of the input. Max pooling was chosen because it was used in AlexNet [4], which was a large influence on this network.
- **Local Response Normalization:** This restricts the values that are possible to a

smaller range while still having the relative ratios between values in the input. This restriction helps reduce overfitting in the model where the model is tuned so well to the training data that it fails to generalize to the test data. This overfitting manifests as a low training error and a high test error. Default parameters from tensorflow are used. Local response normalization was chosen, like max pooling, due to its use in AlexNet [4].

- **Dropout:** When applied to an input layer it randomly zeros out some values according to some probability. The fact that any given input may not be present means that the network cannot depend on any one input feature. This forces the system to learn redundant representations which leads to more general solutions. Dropout was shown to be helpful by Goodfellow *et al.* in their paper [3].
- **Softmax:** Softmax is an algorithm that normalizes the output values of logistic regression. This makes it so the output sums to one and is a valid probability distribution that can be used to make predictions. Softmax is a classic algorithm used in image recognition to turn the results of logistic regression into a probability distribution and so it is used here.
- **Cross Entropy:** Cross entropy is a metric that is used to tell how wrong a prediction is compared to the correct answer. This loss function summarizes how wrong the the model was when comparing its predictions to the real answers. It is the cost function that is used in logistic regression and is therefore used here.
- **AdagradOptimizer:** Is an advanced optimization function that is implemented in tensorflow that takes as inputs the gradients (partial derivatives in respect to each weight

variable) of the cost function and updates these weights. This is an efficient and built into tensorflow so it is a good choice to use.

The deep convolutional neural network is built with four convolutional layers that will transform the input which is size 50 by 50 by 1 (a grayscale image) into a single feature vector of length 128. This feature vector will be passed to six logistic regression classifiers that each have hidden layers. These six classifiers will output the length and the values of the digits in the images. This network will work directly on the pixel values of the input images

1. The first convolutional layer uses a kernel of size 3 by 3 by 1 by 16 and a stride of 1. The layer uses valid padding. This convolution reduces the size of the image from size 50 by 50 by 1 to 48 by 48 by 16. Then max pooling is applied to the input with a stride of two. This results in an image that is size 24 by 24 by 16. This pooling is followed by local response normalization to help reduce overfitting.
2. The second convolutional layer uses a kernel of 1 by 1 by 16 by 32 and a stride of 1. This means that the result of the convolution is the same size as the input but it is much deeper (32 compared to 16). This technique of a convolutional layer that does not reduce the input size but increases the depth is taken from the implementation of a similar network by Goodfellow *et al.* [3]. After this second convolution local response normalization is applied again. Changing this order of pooling and normalization is adapted from the AlexNet architecture used by Krizhevsky *et al.* to win the ImageNet competition in 2012 [4]. Then pooling is applied resulting in an input of 12 by 12 by 32.
3. The third convolution is then applied with a kernel of size 5 by 5 by 32 by 64 and a stride

of 1 to create an output with size 8 by 8 by 64. Pooling is used to change the size to 4 by 4 by 64. Normalization is then applied.

4. The final convolution with kernel 1 by 1 by 64 by 128 and stride 1 is then applied to create a input of 1 by 1 by 128. This input is normalized and then reshaped into a feature vector of size 128.

This final feature vector is then used as input for six different softmax classifiers that have hidden layers of size 16. The first softmax classifier has an output space of 0 to 6 (the length of the sequence in the image) and the rest have output spaces of 0 through 10 (the values of the digits in the image). These softmax classifiers show the probability that a certain digit is that value. by taking the most probable value we can make predictions of the digits in the image. In this network dropout is applied to each layer save input and output.

The results are compared to the labels using “sparse_softmax_cross_entropy_with_logits” to calculate the loss. This cost is then minimized by the AdamOptimizer in order to train the model.

2.4. Benchmark

According to the paper by Goodfellow *et al.* [3] human operators have about 98% accuracy when it comes to identifying multidigit sequences in natural scenes. This would obviously be a good benchmark to try to reach with this system. However the system created by Goodfellow *et al.* [3] only reached 96% accuracy in natural scenes. This state of the art result seems to be a better benchmark than humans levels (it seems more obtainable than trying to beat state of the art results). However this result seems unobtainable due to the development environment. Goodfellow *et al.* [3] achieved 96% using a neural network that was eleven layers deep before being connected to the six final output layers. The final feature vector created by Goodfellow *et al.*’s convolutional network was of size 4096 while the vector created

by this convolutional network is only of size 128 [3]. Goodfellow *et al.* [3] used a distributed framework called DistBelief to train their network and it still took six days to train the network. Time and resources to create such a deep network are lacking. Using a shallower network (about four or five layers) should be able to reach about 90% accuracy according to the graph in Figure 4 of the Goodfellow *et al.* paper [3].

Due to the large amount of samples that have a small length of sequence a poor model might be tempted to guess all 10 (the “not there value”). If the model also guessed all the sequences were length 2 (the mode of the training set) then this model would achieve an accuracy of 60.79%. This sets a floor that the model should try to beat.

3. Methodology

3.1. Data Processing

Data processing for this problem was based on the processing methods described by Goodfellow *et al.* in their paper [3]. First a bounding box that encompasses all the bounding boxes for the digits in the images is found. This box is then scaled up by 30%. The image is then cropped to the size of this scaled up box. This cropped image is then resized to be 50 pixels by 50 pixels. In the paper by Goodfellow *et al.* [3] these images are resized to 64 pixels by 64 pixels. This was the first size tried but the development environment ran out of memory when they were scaled to 64 pixels by 64 pixels, after trying several sizes 50 seemed to be the largest that my development environment could handle. After the resizing the images were converted from the RGB with 3 color channels to grayscale so that rather than three channels the resulting images had only one color channel. Finally the images are normalized by subtracting the mean pixel value from each pixel and dividing it by the standard deviation. This regularization was the only form applied to the dataset. Things like LeCun Local Contrast Normalization were not used.




Figure 3. The same image from Figure 1 after it has been processed (cropped, resized to 50 x 50 pixels and converted to grayscale).

As explained above the one image with length six could be considered an outlier but was not removed from dataset due to high likelihood of similar images being found in the wild

The SVHN dataset is divided into three datasets. The train, test, and extra datasets. The extra dataset is a large collection of easy samples while the train dataset is a smaller dataset with “more difficult” samples. The difficulty assessment for these datasets comes from both a paper by LeCun *et al.* [6] and the sentiment is echoed by the website where the dataset can be downloaded. To create training and validation datasets from the train and extra datasets a method from LeCun *et al.* [6] was used. The Validation set is created so that it is $\frac{2}{3}$ train samples and $\frac{1}{3}$ extra samples from each class. This breaks down into about 4,000 train samples of each class and about 2,000 extra samples from each class (5684 total). Here “class” is defined to be based on what digit is the first in the sequence in the image.

Figure 3 shows Figure 1 after the image has been processed.

3.2. Implementation

The first step in implementation was to obtain the datasets. These were fetched from <http://ufldl.stanford.edu/housenumbers/> using a modified version of the download code from the deep learning course at Udacity <https://www.udacity.com/course/deep-learning--ud730>. The code can be found here <https://github.com/tensorflow/tensorflow/tree/master/tensorflow/examples/udacity>. The data was then extracted using more code modified from the deep learning course.

While implementing the data preprocessing familiarity with numpy was needed to help shape and normalize image arrays. The very tricky part came from the format of the SVHN label file. The labels were saved as a .m matlab file. With the newest versions of matlab python cannot turn the data from a .m file into a dictionary automatically.

The use of the h5py module was new during this project. h5py was used to read data from the .m file and parse it into python dictionaries. This was one of the trickier parts of the project and help from the Udacity forms really helped get a handle on the way that h5py uses references and how to fetch data from it. The code used is adapted from https://github.com/hangyao/street_view_house_numbers/blob/master/3_preprocess_multi.ipynb

Once the preprocessing code was written the next problem that was ran into is the limitations of the development environment. This was all done on a laptop and when the preprocessing code tried to resize all images to be 64 by 64 pixels the program would run out of memory. This meant that a new size must be found. After trying smaller and smaller resize dimensions 50 seemed to be the largest dimensions that the images can be scaled down to without running out of memory.

The Neural Network itself was implemented in

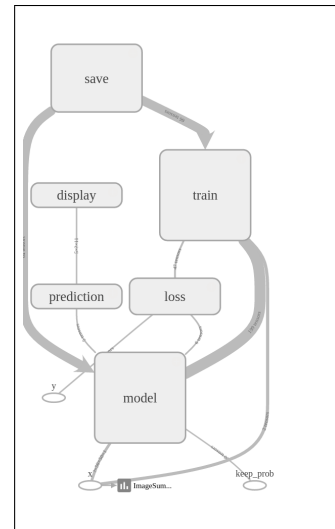


Figure 4. A picture of the tensorflow graph. The model is the network itself and the rest of it is the what allows for training and saving the graph.

tensorflow by Google [1] This was the first time using tensorflow but it was fairly easy. First it was implemented as a straight block of code in an ipython notebook. This was obviously non-conducive to using in an application. The network was rewritten using functions. This took a long time to figure out how to do because of inexperience with tensorflow, especially the fact that things are not evaluated until a session is run, plus the idea of a graph and how to use placeholders took a bit to get used to. This code still looked unreadable so variable scope was added to clean it up. Tensorboard summary operations were also added to help visualize both the model and the learning process. Figure 4 shows the visualization of the entire tensorflow graph. Figure 5 shows the graph of actual network itself.

The trickiest part of the implementation was how to calculate the loss function for this network. Most softmax implementations use cross entropy with the label “one-hot encoded” (the label is a vector that has a 1 in the index of the answer rather than just being a scalar that is the answer). Not wanting to convert labels into “one-hot encoded” labels

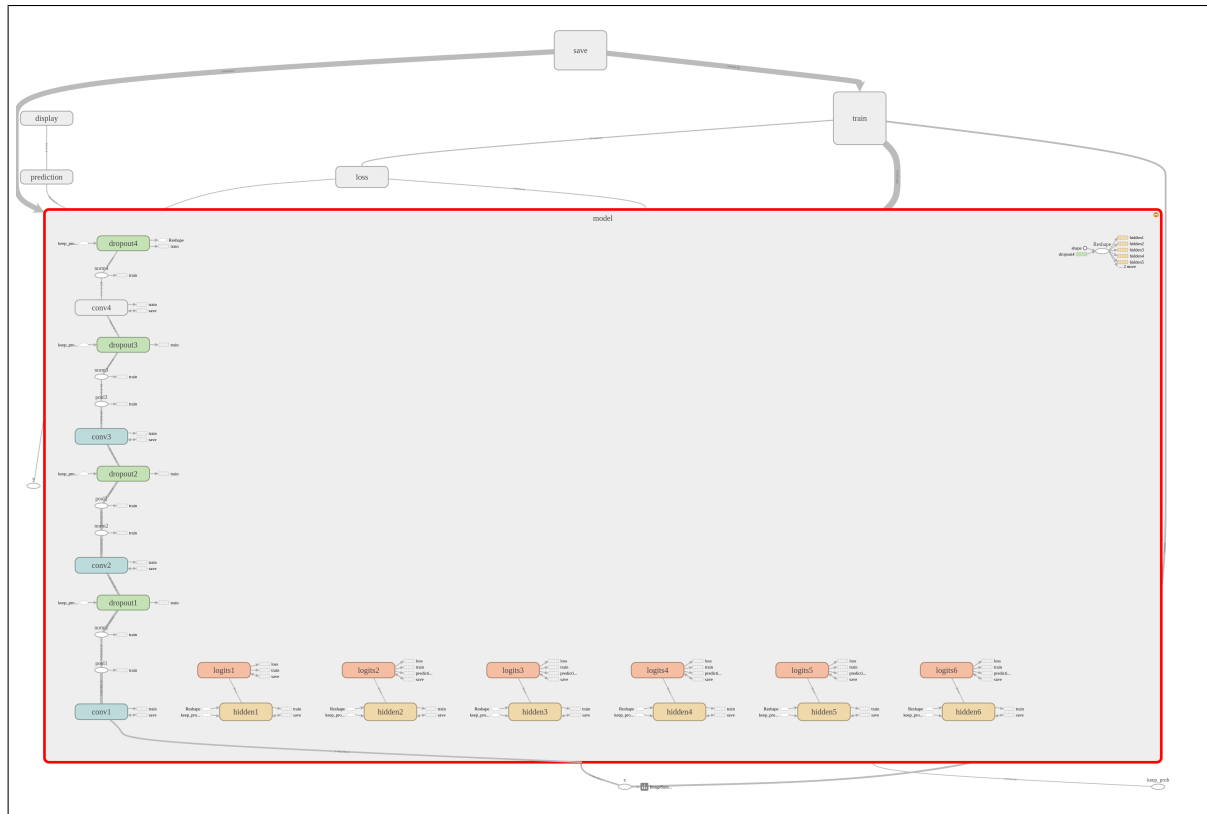


Figure 5. Picture of the network from tensorboard. Larger image available at <http://imgur.com/cx9DINa>

“sparse_softmax_cross_entropy_with_logits”, the tensorflow function that calculates the cross entropy, was used. The use of this function meant that cross entropy could be calculated without needing to use “one-hot encoded” labels (which would have been a drain on system resources). The tensorflow documentation mentions that input to this function should be the output of logistic regression rather than output of softmax because the function does softmax itself to be more efficient. This required a slight rewrite as the inference function that builds network graph originally returned the softmax of the logits layer. The function needed to return the logits layer instead so that it could be used with the “sparse_softmax_cross_entropy_with_logits” function. To calculate the actual loss the result of “sparse_softmax_cross_entropy_with_logits” had to be summed for each label compared to each logits output layer. That is each of




Figure 6. This graph of the loss function over training steps shows how the loss decreased as the model was trained.

the six logits outputs a prediction and the “sparse_softmax_cross_entropy_with_logits” was used to find the cross entropy between each of the logits output and the label that corresponds to that logits. These results were then summed up.

Figure 6 shows the change in the loss function at various training steps. The loss function

trended down during training showing that it is a good loss function and minimizing it helped train the model.

3.3. Refinement

The original model had only three convolutional layers and no hidden layers in the logits. It also only had dropout between the final feature vector and the logits layers. This model had the accuracy reported in the first row of Table 3. While the test accuracy is reported in this table it was not used to tune the model as this would cause the features of the test set to bleed into the training process. The fact that the the training and validation accuracy is low means that the model is not complex enough to represent the data (the model has high bias). Due to these errors the network was made deeper, another convolutional layer was added, and larger, hidden layers were added to the logits.

Once the network was expanded the accuracies were the ones found in the second row of Table 3. The fact that the validation error is so much higher than the test set error (shown by the fact the the accuracy is lower) means that the model is overfitting to the training data. The model has high variance and is having trouble generalizing to the test data and will likely have trouble on other novel data. A common technique to fight overfitting in neural networks is dropout. This was then added to each layer of the model outside of the input and output layers.

Once the network was made larger and dropout was added the accuracy reported in row 3 of Table 3 were found. Similar steps (especially increases in depth) could be taken from here to get even better performance.

4. Results

4.1. Model Evaluation and Validation

The final model is a deep convolutional neural network. This model was chosen due to its use by Goodfellow at all [3] to great success. The

Accuracy	Train	Valid	Test
Original	90.13	88.38	89.45
Second	95.21	90.39	91.11
Best	95.63	92.73	93.89

Table 3. Accuracy of the model on each dataset where the Train dataset is the last accuracy on the last batch of training data.

final model however was deeper than the original plan for the model. This more complex model resulted in very long training times, 11.18 hours on an Intel i5 CPU.

Figure 7 shows the accuracy values for the training and validation datasets as training went on. The validation set was still increasing as the training went on meaning that more training could increase the accuracy. Table 3 includes the final Test accuracy of 93.89%. This accuracy is close to the training and validation accuracy which means that this model does not suffer from high variance. The means that the model does not overfit to the training data and is able to generalize to new data. This model is therefore a good fit to solve this problem.

Outside of using the Test set to test the model the performance was vetted using both a camera application and the ability to run the model on various images. The model preforms reasonably well on these applications. However when these input methods have a slight problem in that they preform rather poor when the digits are not centered in the image. They also preform better when the digits are about 30% smaller than the boundaries of the images. These input methods show that the model is sensitive to slight changes in the input. This could be fixed with a deeper network and more input where the digits are in various places of the input. This problem could also be solved with more convolutional layers because convolutions squeeze the spacial dimensions from the images. Without having the digits in various places in the input images the model never really learns to localize the digits.




Figure 7. A graph showing how the accuracy on the training batch and the accuracy on the validation set changed as the model was trained.




Figure 8. An example of some of the output of the model. The images are images that were fed into the model. The Label is the label for that image from the dataset and the Predict is the value of the output when that image is feed into the network.

The results of this model is quite strong and can be trusted in applications where wrong answers are not the worst thing in the world. If the network was being used for mapping data where a wrong address could mean sending a user to the wrong place then the model would have to be trained more.

This model generates digits found in natural scenes. The model creates usable output and is a good solution to the problem.

4.2. Justification

Figure 8 show several example outputs. It includes the image that was given to the network, the label from the dataset and the predicted digits in the image from the model. The model missed only a single digit from the second image.

This model sufficiently solves the problem. A network that guessed only 2 for the sequence length and only 10 for the each digit would score 60.79%. This is the minimum accuracy required for the system to do better than educated guesses. The test accuracy for the final network is 93.89%. This is far better than the minimum benchmark. 93.89% is also pretty close to the accuracy found by Goodfellow *et al.* [3]. Being close to this benchmark is a very good result.

This network is a good solution to the problem. It has a few selectivity problems as discussed in the “Model evaluation and validation” section. When the improvements mentioned later in the “Improvements” section are applied then these problems may disappear. Even before these improvements the model does a good job classifying digits.

5. Conclusion

5.1. Free Form Visualization

While there is not much more to visualize as the dataset is just images so most of these visualizations will be about the learning process.

Figure 9 shows two images. The left is a processed image that has been resized and converted to grayscale. The right side of the photo is the image after it has passed through one convolutional layer. The changes to the image show how the convolution change the image. It is hard to see the reduced size because it is only two pixels smaller.

The following Figures, 10, 11, and 12 are histogram activations for various weights in the convolutional network. Figure 10 shows the activation for the weights in the third convolutional layer of the network. Figure 11 shows the histogram activation for the hidden layer in the 5th




Figure 9. The image on the left is one of the images that has been processed and is ready to be fed into the network. The image on the right is the image after it has been through a single convolution.




Figure 10. This shows the activation of the third convolutional layer weights in the network. The top and bottom light blue lines show that the training doesn't begin to effect this set of weights until almost 40,000 training steps.

logits layer. Figure 12 shows the activation for one of the final logits layers.

These graphs are a little hard to read but they show the distribution of the activation on each layer. The middle line is line that splits the activations so that 50% fall above it and 50% below. The dark lines are the lines where 69% of the activations fall between them, while the light blue lines are where 89% fall between. All three of these graphs show what the majority of activations (between the blue lines) are. The differences between the graph shows how the different layers are trained. The deepest layer (convolution 3 weights in Figure 10) is the slowest to change. This change is visible due to the expansion in the




Figure 11. This shows the activation of the hidden layer weights in the fifth logits in the network. The top and bottom light blue lines show that the training doesn't begin to effect this set of weights until almost 30,000 training steps.



Figure 12. This shows the activation of the output layer weights in the first logits in the network. The top and bottom light blue lines show that the training starts to change weights almost immediately.

y dimension of the top and bottom lines in the graph. The hidden layer changes quicker and the logits changes the quickest in that it already starts with a wide distribution and stays wide. The explanation for the difference between these graphs is explained by the saturation of the logistic function at different levels of the neural network [2]. These graphs help visualize the learning process and in future runs and improvements these graphs will help see patterns and show that the changes are helping the models performance.

5.2. Reflection

The first step of this project was to read the existing literature available on the topic. The most helpful papers were by Goodfellow *et al.* [3] and

LeCun *et al.* [6]. The Goodfellow *et al.* [3] paper was especially helpful. That is where the idea of creating a single convolutional feature vector and using that as input to various logistic classifiers.

The next step was to analyze and process the dataset. This was an interesting part of the process because the use of h5py was new and without this project it would probably not have been used. The h5py module is used to access data structures that are saved on to disk.

After processing the data the model was build in tensorflow. This was an interesting part to the project because it was new. Figuring out how to build the convolutional network so that the final feature vector was only size 1 by 1 by depth. The difficult part was turning the mess of spaghetti code into a well organized function, several errors were introduced in the process that resulted in very low accuracy (the cross entropy error did take into account the 4th logits). Figuring out the cause of this error took several days and was the hardest part of the project.

The final steps was adding the network into an application. This was interesting to use opencv for the first time. After hooking all the pieces together the last thing to do was train the model. This took 11 hours and the anticipation was one of the harder parts of the project.

This is one of the largest end-to-end projects that have been undertaken and while the results are a little weaker than the goal was, the project was a large success.

5.3. Improvement

There is a lot of improvement that can be made from here. One of the first things that might be able to help performance would be to normalize the data after it is split into the training and validation set rather than before. This would help make the images within a dataset more uniform.

Another problem that the model has is that it struggles when the digit sequence is not in the center of the image. This problem, plus the problem that the accuracy is not as high as it could

be, can be solved in a single tweak. As discussed in the paper by Goodfellow *et al.* [3] more data can be artificially created by grabbing different parts of the bounding box. For example once you blow up the largest bounding box you can grab the whole thing as one data point, you can grab top left box (from the scaled up corner top left to the original bottom right corner) so the image is in the lower right part of the image and so on. This means that a single image can be used to create 5 training points. This would help train the model to detect digits where ever they appear in the image. The increase in amount of data will also help the accuracy of the model.

In order to do some of the other improvements a better development environment is needed. One obvious improvement is during data preprocessing resize the images to 64 by 64 rather than 50 by 50. This allows for more pixels and clearer images. This will help performance but requires more memory for the computer.

As the Figure 6 loss function graph shows that at the end of the training steps the loss function was still decreasing. The same trend is shown in Figure 7 as the accuracy continues to increase. This means that the model would benefit from more training. The training time for the best model was 11 hours long so in order to realistically train for significantly more steps would require a GPU to speed up training and allow for more training.

Another improvement is to create a more complex model. This means a deeper convolutional network and more hidden units in the logits layers as well as more hidden layers. This also can mean using larger convolutional kernels and same padding. A more complex network also means that it is possible to over fit so dropout is needed. For this deeper network to be feasible parallel training on a GPU would be needed for training. The largest indicator that the deeper network would help is the fact that the Goodfellow *et al.* [3] paper had a feature vector produced by the convolutional part of the network had a 4096 fea-

tures compared to the 128 features in this model.

The majority of improvements to this network boils down to having a large network and more data. This requires an improved development environment. This work should be easy to expand on and will be soon.

References

- [1] M. Abadi, A. Agarwal, P. Barham, E. Brevdo, Z. Chen, C. Citro, G. S. Corrado, A. Davis, J. Dean, M. Devin, S. Ghemawat, I. Goodfellow, A. Harp, G. Irving, M. Isard, Y. Jia, R. Jozefowicz, L. Kaiser, M. Kudlur, J. Levenberg, D. Mané, R. Monga, S. Moore, D. Murray, C. Olah, M. Schuster, J. Shlens, B. Steiner, I. Sutskever, K. Talwar, P. Tucker, V. Vanhoucke, V. Vasudevan, F. Viégas, O. Vinyals, P. Warden, M. Wattenberg, M. Wicke, Y. Yu, and X. Zheng. TensorFlow: Large-scale machine learning on heterogeneous systems, 2015. Software available from tensorflow.org.
- [2] X. Glorot and Y. Bengio. Understanding the difficulty of training deep feedforward neural networks, 2010.
- [3] I. J. Goodfellow, Y. Bulatov, J. Ibarz, S. Arnoud, and V. Shet. Multi-digit number recognition from street view imagery using deep convolutional neural networks, 2014.
- [4] A. Krizhevsky, I. Sutskever, and G. E. Hinton. Imagenet classification with deep convolutional neural networks, 2012.
- [5] Y. LeCun, L. D. Jackel, B. Boser, J. S. D. amd H. P. Graf, I. Guyon, D. Henderson, R. E. Howard, and W. Hubbard. Handwritten digit recognition: Applications of neural net chips and automatic learning, 1999.
- [6] P. Sermanet, S. Chintala, and Y. LeCun. Convolutional neural networks applied to house numbers digit classification, 2012.